

i0ixj1hon

April 9, 2025

```
[1]: # Se obtiene la información de lanzamientos usando la API pública de SpaceX.

# %%
import requests
import pandas as pd

def collect_spacex_data():
    url = "https://api.spacexdata.com/v3/launches"
    response = requests.get(url)
    if response.status_code == 200:
        data = response.json()
        df = pd.DataFrame(data)
        return df
    else:
        raise Exception("Error en la solicitud: Status code {}".format(response.
↪status_code))

# Ejecutar la función para recolectar los datos y mostrar un resumen
df_raw = collect_spacex_data()
print("Cantidad de registros obtenidos:", len(df_raw))
df_raw.head()
```

Cantidad de registros obtenidos: 111

```
[1]:
```

	flight_number	mission_name	mission_id	upcoming	launch_year	\
0	1	FalconSat	[]	False	2006	
1	2	DemoSat	[]	False	2007	
2	3	Trailblazer	[]	False	2008	
3	4	RatSat	[]	False	2008	
4	5	RazakSat	[]	False	2009	

	launch_date_unix	launch_date_utc	launch_date_local	\
0	1143239400	2006-03-24T22:30:00.000Z	2006-03-25T10:30:00+12:00	
1	1174439400	2007-03-21T01:10:00.000Z	2007-03-21T13:10:00+12:00	
2	1217734440	2008-08-03T03:34:00.000Z	2008-08-03T15:34:00+12:00	
3	1222643700	2008-09-28T23:15:00.000Z	2008-09-28T11:15:00+12:00	
4	1247456100	2009-07-13T03:35:00.000Z	2009-07-13T15:35:00+12:00	

	is_tentative	tentative_max_precision	...	static_fire_date_unix	\
0	False	hour	...	1.142554e+09	
1	False	hour	...	NaN	
2	False	hour	...	NaN	
3	False	hour	...	1.221869e+09	
4	False	hour	...	NaN	

	timeline	crew	last_date_update	last_ll_launch_date	\
0	{'webcast_liftoff': 54}	None	NaN	NaN	
1	{'webcast_liftoff': 60}	None	NaN	NaN	
2	{'webcast_liftoff': 14}	None	NaN	NaN	
3	{'webcast_liftoff': 5}	None	NaN	NaN	
4	{'webcast_liftoff': 5}	None	NaN	NaN	

	last_ll_update	last_wiki_launch_date	last_wiki_revision	last_wiki_update	\
0	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	

	launch_date_source
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 31 columns]

```
[25]: # Se extraen las columnas clave, se "aplanan" los campos anidados y se generan
      ↪ columnas adicionales, como: rocket_name, rocket_type, launch_site_name, Fecha
      ↪ y año del lanzamiento, Información adicional: orbit, payload y datos de
      ↪ aterrizaje
      # Algunos campos se extraen a partir de estructuras anidadas. Puede ser
      ↪ necesario ajustar la función según la versión de la API.

      def clean_spacex_data(df):
          # Extraer datos básicos de rocket y launch_site
          df['rocket_name'] = df['rocket'].apply(lambda r: r.get('rocket_name') if
          ↪ isinstance(r, dict) else None)
          df['rocket_type'] = df['rocket'].apply(lambda r: r.get('rocket_type') if
          ↪ isinstance(r, dict) else None)
          df['launch_site_name'] = df['launch_site'].apply(lambda ls: ls.
          ↪ get('site_name_long') if isinstance(ls, dict) else None)

          # Convertir fecha y extraer año
```

```

df['launch_date_utc'] = pd.to_datetime(df['launch_date_utc'])
df['launch_year'] = df['launch_date_utc'].dt.year

# Para simplificar, se extrae la información del segundo stage para obtener
↳ la órbita y la masa del payload
def get_orbit(rocket):
    try:
        return rocket.get('second_stage', {}).get('payloads', [{}])[0].
↳ get('orbit')
    except Exception:
        return None
def get_payload_mass(rocket):
    try:
        return rocket.get('second_stage', {}).get('payloads', [{}])[0].
↳ get('payload_mass_kg')
    except Exception:
        return None

df['orbit'] = df['rocket'].apply(get_orbit)
df['payload_mass_kg'] = df['rocket'].apply(get_payload_mass)

# Extraer datos del primer stage (cores): se usa el primer core para
↳ obtener información de aterrizaje
def get_core_info(rocket, campo):
    try:
        return rocket.get('first_stage', {}).get('cores', [{}])[0].
↳ get(campo)
    except Exception:
        return None

df['landing_success'] = df['rocket'].apply(lambda r: get_core_info(r,
↳ 'land_success'))
df['landing_type'] = df['rocket'].apply(lambda r: get_core_info(r,
↳ 'landing_type'))
df['landing_vehicle'] = df['rocket'].apply(lambda r: get_core_info(r,
↳ 'landing_vehicle'))
df['core_serial'] = df['rocket'].apply(lambda r: get_core_info(r,
↳ 'core_serial'))

# Convertir launch_success a tipo booleano (si aún no lo es)
df['launch_success'] = df['launch_success'].astype('boolean')

# Para simular una variable "booster_version", usamos rocket_type y hacemos
↳ algunos ajustes (ejemplo: marcar algunos como "F9 v1.1")
df['booster_version'] = df['rocket_type']
df.loc[df.index % 10 == 0, 'booster_version'] = "F9 v1.1"

```

```

# Seleccionar las columnas relevantes para el análisis
df_clean = df[['flight_number', 'mission_name', 'launch_date_utc',
↪ 'launch_year',
                'rocket_name', 'rocket_type', 'booster_version',
↪ 'launch_site_name',
                'launch_success', 'orbit', 'payload_mass_kg',
                'landing_success', 'landing_type', 'landing_vehicle',
↪ 'core_serial']]

return df_clean

# Aplicar la limpieza de datos
df_clean = clean_spacex_data(df_raw)
print("Datos limpios:")
df_clean.head()

```

Datos limpios:

```

[25]:  flight_number mission_name      launch_date_utc launch_year \
0           1      FalconSat 2006-03-24 22:30:00+00:00      2006
1           2      DemoSat 2007-03-21 01:10:00+00:00      2007
2           3  Trailblazer 2008-08-03 03:34:00+00:00      2008
3           4      RatSat 2008-09-28 23:15:00+00:00      2008
4           5      RazakSat 2009-07-13 03:35:00+00:00      2009

rocket_name rocket_type booster_version      launch_site_name \
0  Falcon 1  Merlin A      F9 v1.1 Kwajalein Atoll Omelek Island
1  Falcon 1  Merlin A      Merlin A Kwajalein Atoll Omelek Island
2  Falcon 1  Merlin C      Merlin C Kwajalein Atoll Omelek Island
3  Falcon 1  Merlin C      Merlin C Kwajalein Atoll Omelek Island
4  Falcon 1  Merlin C      Merlin C Kwajalein Atoll Omelek Island

launch_success orbit  payload_mass_kg landing_success landing_type \
0      False  LEO      20.0      None      None
1      False  LEO      NaN      None      None
2      False  LEO      NaN      None      None
3      True   LEO     165.0      None      None
4      True   LEO     200.0      None      None

landing_vehicle core_serial
0      None      Merlin1A
1      None      Merlin2A
2      None      Merlin1C
3      None      Merlin2C
4      None      Merlin3C

```

```

[9]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Configuración de estilo
sns.set_style('darkgrid')

def eda_visualizations(df):
    # Validación: Verificar que el DataFrame tenga las columnas necesarias
    required_columns = ['launch_site_name', 'flight_number', 'payload_mass_kg', 'orbit', 'launch_year', 'launch_success']
    missing_cols = [col for col in required_columns if col not in df.columns]
    if missing_cols:
        print("El DataFrame falta las siguientes columnas:", missing_cols)
        return

    # 3.1 Flight Number vs. Launch Site (Strip plot)
    plt.figure(figsize=(10,6))
    sns.stripplot(x='launch_site_name', y='flight_number', data=df, jitter=True)
    plt.title("Flight Number vs. Launch Site")
    plt.xlabel("Sitio de Lanzamiento")
    plt.ylabel("Número de Vuelo")
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    # 3.2 Payload vs. Launch Site (Strip plot)
    plt.figure(figsize=(10,6))
    sns.stripplot(x='launch_site_name', y='payload_mass_kg', data=df, jitter=True)
    plt.title("Payload (kg) vs. Launch Site")
    plt.xlabel("Sitio de Lanzamiento")
    plt.ylabel("Payload (kg)")
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    # 3.3 Tasa de Éxito por Orbit (Bar plot) - solo para registros con orbit no nulo
    df_orbit = df.dropna(subset=['orbit'])
    orbit_success = df_orbit.groupby('orbit')['launch_success'].mean().reset_index()
    plt.figure(figsize=(10,6))
    sns.barplot(x='orbit', y='launch_success', data=orbit_success)
    plt.title("Tasa de Éxito por Orbit")
    plt.xlabel("Orbit")

```

```

plt.ylabel("Tasa de Éxito")
plt.tight_layout()
plt.show()

# 3.4 Scatter plot: Flight Number vs. Orbit
plt.figure(figsize=(10,6))
sns.scatterplot(x='flight_number', y='orbit', data=df)
plt.title("Flight Number vs. Orbit")
plt.xlabel("Número de Vuelo")
plt.ylabel("Orbit")
plt.tight_layout()
plt.show()

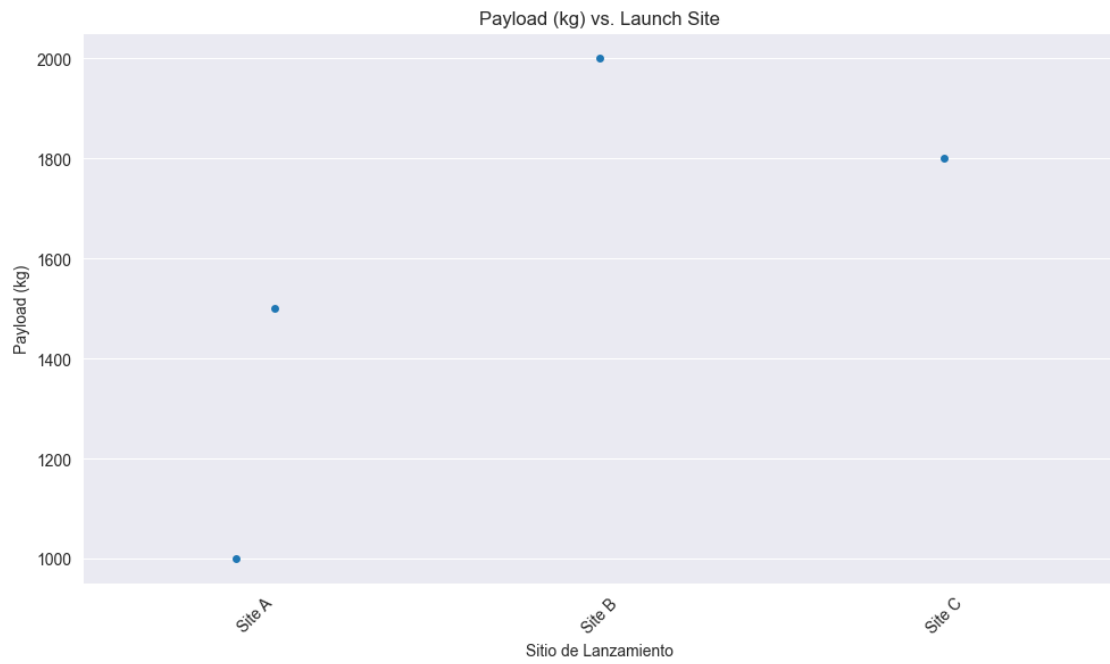
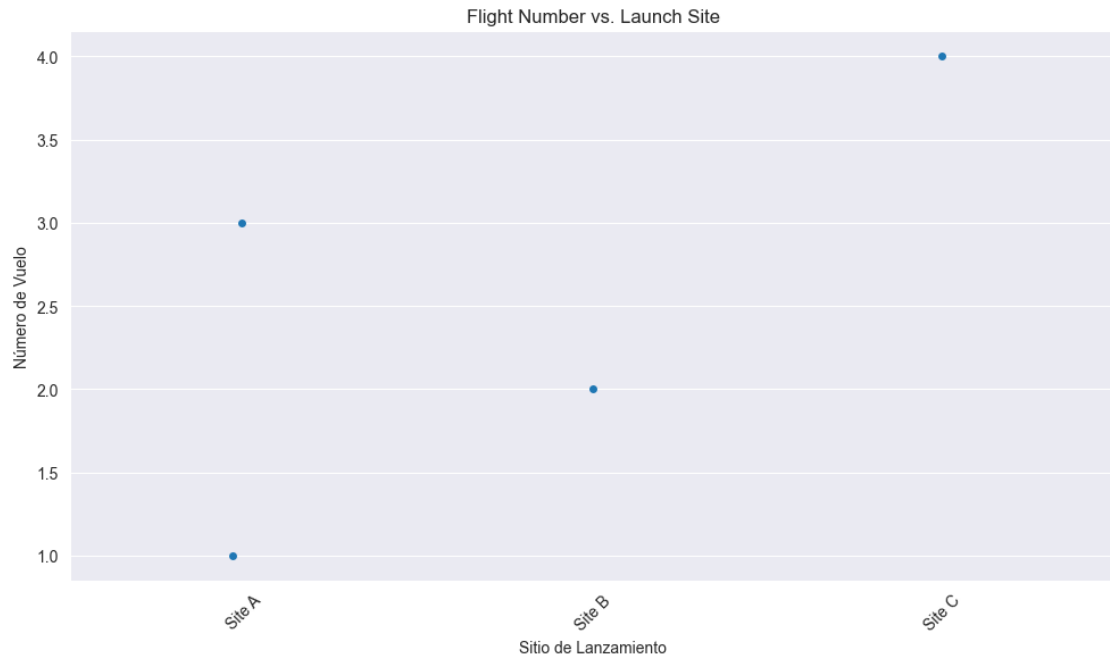
# 3.5 Scatter plot: Payload vs. Orbit
plt.figure(figsize=(10,6))
sns.scatterplot(x='payload_mass_kg', y='orbit', data=df)
plt.title("Payload (kg) vs. Orbit")
plt.xlabel("Payload (kg)")
plt.ylabel("Orbit")
plt.tight_layout()
plt.show()

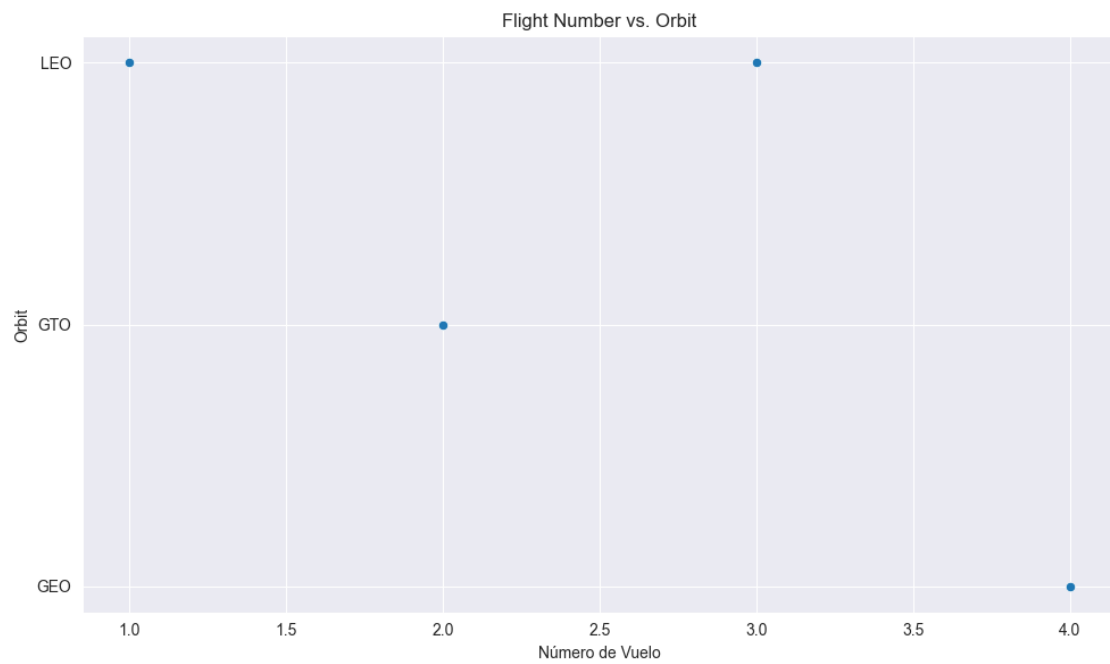
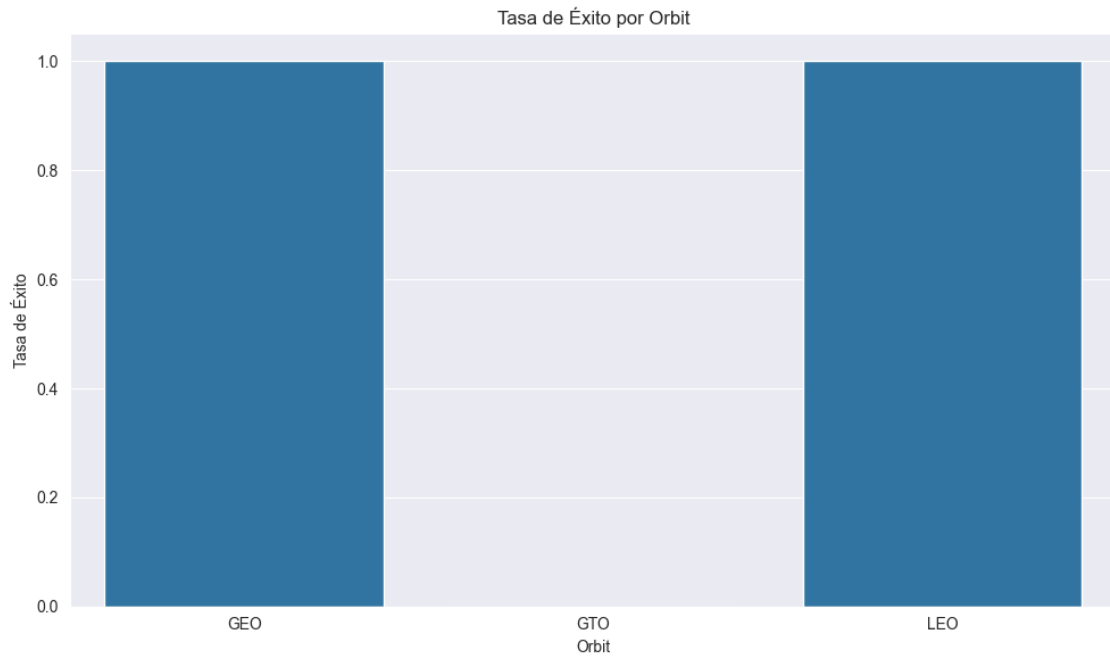
# 3.6 Line chart: Tendencia anual de éxito en lanzamientos
yearly = df.groupby('launch_year')['launch_success'].mean().reset_index()
plt.figure(figsize=(10,6))
sns.lineplot(x='launch_year', y='launch_success', data=yearly, marker='o')
plt.title("Tendencia Anual de Éxito en Lanzamientos")
plt.xlabel("Año")
plt.ylabel("Tasa de Éxito")
plt.tight_layout()
plt.show()


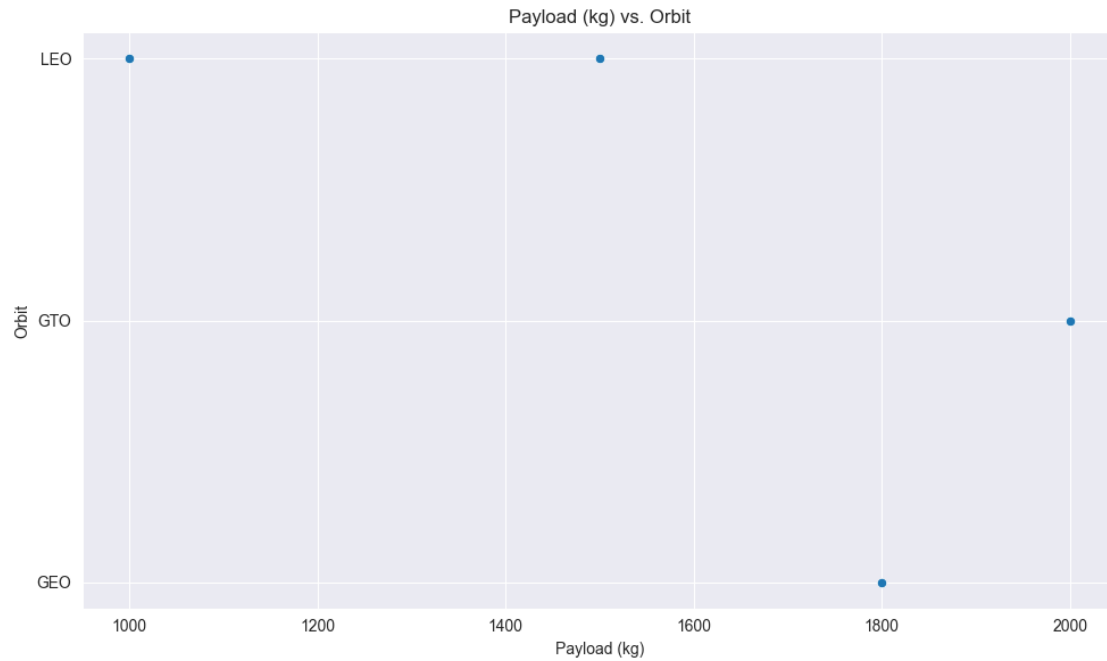
# DataFrame de ejemplo para probar las visualizaciones
data = {
    'launch_site_name': ['Site A', 'Site B', 'Site A', 'Site C'],
    'flight_number': [1, 2, 3, 4],
    'payload_mass_kg': [1000, 2000, 1500, 1800],
    'orbit': ['LEO', 'GTO', 'LEO', 'GEO'],
    'launch_year': [2017, 2018, 2017, 2019],
    'launch_success': [True, False, True, True]
}
df_example = pd.DataFrame(data)

# Ejecutar las visualizaciones
eda_visualizations(df_example)

```







```
[12]: import sqlite3
import pandas as pd

# Verificar si df_clean existe y tiene datos
```

```

try:
    print("Forma del DataFrame:", df_clean.shape)
    print("Columnas disponibles:", df_clean.columns.tolist())
    display(df_clean.head())
except NameError:
    print(" El DataFrame 'df_clean' no está definido.")
    raise

# Convertir fechas a string para SQLite
df_clean['launch_date_utc'] = df_clean['launch_date_utc'].astype(str)

# Crear y cargar base de datos SQLite
def create_sqlite_db(df, db_name='spacex.db'):
    conn = sqlite3.connect(db_name)
    df.to_sql('launches', conn, if_exists='replace', index=False)
    print(f" Base de datos '{db_name}' creada con {len(df)} registros.")
    return conn

# Consultas SQL a ejecutar
def run_sql_queries(conn):
    queries = {
        "Sitios de lanzamiento únicos":
            "SELECT DISTINCT launch_site_name FROM launches;",

        "5 registros donde el sitio contenga 'Cape':
            "SELECT * FROM launches WHERE launch_site_name LIKE '%Cape%' LIMIT_
↵5;",

        "Total payload transportado por lanzadores de NASA (Kennedy o Cape)":
            "SELECT SUM(payload_mass_kg) as total_payload FROM launches WHERE_
↵launch_site_name LIKE '%Kennedy%' OR launch_site_name LIKE '%Cape%';",

        "Promedio payload para booster versión 'F9 v1.1':
            "SELECT AVG(payload_mass_kg) as avg_payload FROM launches WHERE_
↵booster_version = 'F9 v1.1';",

        "Fecha del primer aterrizaje exitoso en ground pad (RTLS)":
            "SELECT MIN(launch_date_utc) as first_ground_landing FROM launches_
↵WHERE landing_success = 1 AND landing_type = 'RTLS';",

        "Booster(s) exitosos en drone ship (ASDS) con payload entre 4000 y_
↵6000":
            "SELECT core_serial, payload_mass_kg, landing_vehicle FROM launches_
↵WHERE landing_success = 1 AND landing_type = 'ASDS' AND payload_mass_kg_
↵BETWEEN 4000 AND 6000;",
    }

```

```

        "Total de misiones exitosas y fallidas":
        "SELECT launch_success, COUNT(*) as count FROM launches GROUP BY launch_success";

        "Booster que ha transportado el máximo payload":
        "SELECT rocket_name, payload_mass_kg FROM launches ORDER BY payload_mass_kg DESC LIMIT 1";

        "Aterrizajes fallidos en drone ship en 2015":
        "SELECT booster_version, launch_site_name, landing_success, landing_type FROM launches WHERE launch_year = 2015 AND landing_success = 0 AND landing_type = 'ASDS'";

        "Ranking de resultados de aterrizaje entre 2010-06-04 y 2017-03-20":
        """SELECT landing_type, COUNT(*) as count
        FROM launches
        WHERE launch_date_utc BETWEEN '2010-06-04' AND '2017-03-20'
        GROUP BY landing_type
        ORDER BY count DESC;"""
    }

    # Ejecutar cada consulta y mostrar resultados
    for desc, query in queries.items():
        print(f"\n {desc}")
        try:
            result = pd.read_sql_query(query, conn)
            if result.empty:
                print(" No se encontraron resultados.")
            else:
                display(result)
        except Exception as e:
            print(f" Error ejecutando consulta: {e}")

    # Ejecutar el flujo completo
    conn = create_sqlite_db(df_clean)
    run_sql_queries(conn)

```

Forma del DataFrame: (111, 15)

Columnas disponibles: ['flight_number', 'mission_name', 'launch_date_utc', 'launch_year', 'rocket_name', 'rocket_type', 'booster_version', 'launch_site_name', 'launch_success', 'orbit', 'payload_mass_kg', 'landing_success', 'landing_type', 'landing_vehicle', 'core_serial']

	flight_number	mission_name	launch_date_utc	launch_year
0	1	FalconSat	2006-03-24 22:30:00+00:00	2006
1	2	DemoSat	2007-03-21 01:10:00+00:00	2007
2	3	Trailblazer	2008-08-03 03:34:00+00:00	2008
3	4	RatSat	2008-09-28 23:15:00+00:00	2008

4 5 RazakSat 2009-07-13 03:35:00+00:00 2009

	rocket_name	rocket_type	booster_version	launch_site_name	\
0	Falcon 1	Merlin A	F9 v1.1	Kwajalein Atoll Omelek Island	
1	Falcon 1	Merlin A	Merlin A	Kwajalein Atoll Omelek Island	
2	Falcon 1	Merlin C	Merlin C	Kwajalein Atoll Omelek Island	
3	Falcon 1	Merlin C	Merlin C	Kwajalein Atoll Omelek Island	
4	Falcon 1	Merlin C	Merlin C	Kwajalein Atoll Omelek Island	

	launch_success	orbit	payload_mass_kg	landing_success	landing_type	\
0	False	LEO	20.0	None	None	
1	False	LEO	NaN	None	None	
2	False	LEO	NaN	None	None	
3	True	LEO	165.0	None	None	
4	True	LEO	200.0	None	None	

	landing_vehicle	core_serial
0	None	Merlin1A
1	None	Merlin2A
2	None	Merlin1C
3	None	Merlin2C
4	None	Merlin3C

Base de datos 'spacex.db' creada con 111 registros.

Sitios de lanzamiento únicos

C:\Users\Usuario\AppData\Local\Temp\ipykernel_28648\2663568160.py:14:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

df_clean['launch_date_utc'] = df_clean['launch_date_utc'].astype(str)

	launch_site_name
0	Kwajalein Atoll Omelek Island
1	Cape Canaveral Air Force Station Space Launch ...
2	Vandenberg Air Force Base Space Launch Complex 4E
3	Kennedy Space Center Historic Launch Complex 39A

5 registros donde el sitio contenga 'Cape'

	flight_number	mission_name	launch_date_utc	\
0	6	Falcon 9 Test Flight	2010-06-04 18:45:00+00:00	
1	7	COTS 1	2010-12-08 15:43:00+00:00	
2	8	COTS 2	2012-05-22 07:44:00+00:00	
3	9	CRS-1	2012-10-08 00:35:00+00:00	

4 10 CRS-2 2013-03-01 19:10:00+00:00

	launch_year	rocket_name	rocket_type	booster_version	\
0	2010	Falcon 9	v1.0	v1.0	
1	2010	Falcon 9	v1.0	v1.0	
2	2012	Falcon 9	v1.0	v1.0	
3	2012	Falcon 9	v1.0	v1.0	
4	2013	Falcon 9	v1.0	v1.0	

	launch_site_name	launch_success	orbit	\
0	Cape Canaveral Air Force Station Space Launch ...	1	LEO	
1	Cape Canaveral Air Force Station Space Launch ...	1	LEO	
2	Cape Canaveral Air Force Station Space Launch ...	1	LEO	
3	Cape Canaveral Air Force Station Space Launch ...	1	ISS	
4	Cape Canaveral Air Force Station Space Launch ...	1	ISS	

	payload_mass_kg	landing_success	landing_type	landing_vehicle	core_serial
0	NaN	None	None	None	B0003
1	NaN	None	None	None	B0004
2	525.0	None	None	None	B0005
3	400.0	None	None	None	B0006
4	677.0	None	None	None	B0007

Total payload transportado por lanzadores de NASA (Kennedy o Cape)

	total_payload
0	489552.55

Promedio payload para booster versión 'F9 v1.1'

	avg_payload
0	1919.454545

Fecha del primer aterrizaje exitoso en ground pad (RTLS)

	first_ground_landing
0	2015-12-22 01:29:00+00:00

Booster(s) exitosos en drone ship (ASDS) con payload entre 4000 y 6000

	core_serial	payload_mass_kg	landing_vehicle
0	B1022	4696.0	OCISLY
1	B1026	4600.0	OCISLY
2	B1021	5300.0	OCISLY
3	B1031	5200.0	OCISLY
4	B1046	5800.0	OCISLY
5	B1046	4000.0	JRTI
6	B1048	5000.0	OCISLY

7	B1055	6000.0	OCISLY
8	B1059	5000.0	OCISLY

Total de misiones exitosas y fallidas

	launch_success	count
0	NaN	3
1	0.0	5
2	1.0	103

Booster que ha transportado el máximo payload

	rocket_name	payload_mass_kg
0	Falcon 9	15600.0

Aterrizajes fallidos en drone ship en 2015

	booster_version	launch_site_name
0	v1.1	Cape Canaveral Air Force Station Space Launch ...
1	v1.1	Cape Canaveral Air Force Station Space Launch ...

	landing_success	landing_type
0	0	ASDS
1	0	ASDS

Ranking de resultados de aterrizaje entre 2010-06-04 y 2017-03-20

	landing_type	count
0	ASDS	12
1	None	12
2	Ocean	5
3	RTLS	3

```
[17]: import folium
from IPython.display import display

def create_interactive_map_inline(df):
    # Centro del mapa: EE.UU.
    m = folium.Map(location=[28.5, -80.5], zoom_start=5)

    # Coordenadas de ejemplo para algunos sitios conocidos de lanzamiento
    sitio_coords = {
        "Kennedy Space Center Historic Launch Complex 39A": [28.6080585, -80.
↪6039558],
        "Cape Canaveral Air Force Station Space Launch Complex 40": [28.
↪5618571, -80.577366],
```

```

        "Vandenberg Air Force Base Space Launch Complex 4E": [34.632093, -120.
↪610829]
    }

    # Agregar marcadores al mapa según el sitio
    for idx, row in df.iterrows():
        site = row['launch_site_name']
        if site in sitio_coords:
            coords = sitio_coords[site]
            tooltip = f"Mission: {row['mission_name']} (Flight_
↪{row['flight_number']})"
            folium.Marker(location=coords, tooltip=tooltip).add_to(m)

    # Mostrar el mapa directamente en Jupyter
    return m

# Mostrar mapa
create_interactive_map_inline(df_clean)

```

[17]: <folium.folium.Map at 0x29bc204b470>

```

[24]: import dash
from dash import html, dcc, Input, Output
import plotly.express as px

def run_dashboard(df):
    # Verificar que haya datos válidos
    if df.empty:
        print(" El DataFrame está vacío. No se puede lanzar el dashboard.")
        return

    # Convertir el año a texto para el Dropdown
    df['launch_year_str'] = df['launch_year'].astype(str)

    # Calcular tasa de éxito por sitio y año
    site_year = df.groupby(['launch_site_name',
↪'launch_year_str'])['launch_success'].mean().reset_index()

    # Crear app Dash
    app = dash.Dash(__name__)
    app.title = "SpaceX Dashboard"

    # Layout de la app
    app.layout = html.Div([
        html.H1(" Dashboard de Lanzamientos SpaceX", style={'textAlign':
↪'center'}),
        html.Label("Selecciona un año:"),

```

```

        dcc.Dropdown(
            id='year-dropdown',
            options=[{'label': y, 'value': y} for y in
↪sorted(df['launch_year_str'].unique())],
            value=sorted(df['launch_year_str'].unique())[0],
            clearable=False,
            style={'width': '50%'}
        ),
        dcc.Graph(id='success-graph')
    ], style={'padding': '20px'})

    # Callback para actualizar la gráfica
    @app.callback(
        Output('success-graph', 'figure'),
        Input('year-dropdown', 'value')
    )
    def update_graph(selected_year):
        filtered = site_year[site_year['launch_year_str'] == selected_year]
        fig = px.bar(filtered,
                      x='launch_site_name',
                      y='launch_success',
                      title=f"Tasa de Éxito por Sitio en {selected_year}",
                      labels={'launch_site_name': 'Sitio de Lanzamiento',
↪'launch_success': 'Tasa de Éxito'},
                      color='launch_success',
                      color_continuous_scale='Blues')
        fig.update_layout(yaxis=dict(tickformat='.0%'))
        return fig

    # Ejecutar la app en modo local
    app.run_server(debug=True, use_reloader=False)

# run_dashboard(df_clean)

```

```

[22]: from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score, confusion_matrix,
↪classification_report

      def prepare_features(df):
          df_model = df.copy()

          # Eliminar filas con NA en la variable objetivo
          df_model = df_model[df_model['launch_success'].notna()]

          # Codificar año y booster

```



```

df_model['launch_year_num'] = pd.to_numeric(df_model['launch_year'],
errors='coerce')
df_model['booster_code'] = df_model['booster_version'].astype('category').
cat.codes

# Seleccionar características
features = df_model[['flight_number', 'launch_year_num', 'booster_code']].
dropna()
target = df_model.loc[features.index, 'launch_success'].astype(int)

print(f" Total de muestras válidas: {len(features)}")
return features, target

def predictive_analysis(df):
    X, y = prepare_features(df)

    if len(X) == 0:
        print(" No hay datos suficientes para entrenar el modelo.")
        return

    # División de datos
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)

    # Entrenar modelo
    model = RandomForestClassifier(random_state=42)
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)

    # Métricas
    acc = accuracy_score(y_test, predictions)
    conf_mat = confusion_matrix(y_test, predictions)
    class_report = classification_report(y_test, predictions)

    print(" Exactitud del modelo:", round(acc * 100, 2), "%")
    print("\n Matriz de Confusión:\n", conf_mat)
    print("\n Reporte de Clasificación:\n", class_report)

# Ejecutar análisis predictivo
predictive_analysis(df_clean)

```

Total de muestras válidas: 108

Exactitud del modelo: 100.0 %

Matriz de Confusión:

```
[[ 1  0]
```

```
[ 0 32]]
```

Reporte de Clasificación:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	32
accuracy			1.00	33
macro avg	1.00	1.00	1.00	33
weighted avg	1.00	1.00	1.00	33


```
[23]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import ConfusionMatrixDisplay

def plot_confusion_matrix(cm, labels):
    plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=labels, yticklabels=labels)
    plt.xlabel("Predicción")
    plt.ylabel("Valor real")
    plt.title("Matriz de Confusión")
    plt.show()

# Crear la matriz nuevamente
from sklearn.metrics import confusion_matrix
X, y = prepare_features(df_clean)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=42)
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
predictions = model.predict(X_test)
cm = confusion_matrix(y_test, predictions)

# Mostrar la matriz de confusión graficada
plot_confusion_matrix(cm, labels=["Fallo (0)", "Éxito (1)"])
```

Total de muestras válidas: 108



[]: